


A Framework for Double-Blind Federated Adaptation of Foundation Models

Background



Double-Blind Constraints

Model Privacy: Clients cannot access the FM.

Data Privacy: The LSP cannot access the local data.

Goal

Train a lightweight **parallel adapter** and a **classification head** to maximize performance while preserving privacy.

Method


BlindFed: Recipe Card

Assumptions. Public auxiliary $\mathcal{D}_{aux}$; modular FM with L blocks; thin clients, powerful server; semi-honest parties.

Protocol.

- Pre-processing:** Server: Convert FM to FHE-friendly form (approximate Softmax, GELU, LayerNorm) and distill from original FM on $\mathcal{D}_{aux}$. Clients: Send FHE encrypted data to the server, only once.
- Per-round:**
 - Server performs encrypted block-wise inference and permutes IR (intermediate repr.) batches.
 - Clients decrypt IRs and train parallel adapter + head locally (no FM exposure).
 - Server aggregates updates via MPC; optional SBS.

Double-Blind Adaptation of Foundation Models




Nurbek Tastan, MBZUAI

Karthik Nandakumar, MBZUAI, MSU




Experiments

Main Results





BlindFed significantly outperforms DB-LP, while achieving accuracy only marginally lower than non-DB methods.

Scaling with Number of Clients



BlindFed achieves strong performance across different no. of clients with far lower latency and parameter count.

OOD Auxiliary Data & Overhead



Mohamed bin Zayed
University of
Artificial Intelligence

MICHIGAN STATE
UNIVERSITY